

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from tabulate import tabulate
from scipy.stats import multivariate_normal as mvn

# Model classes

# Naive Bayes - class
class GaussNB():
    def fit(self, X, y, epsilon=1e-3):

        self.likelihoods = dict ()
        self.priors = dict()

        self.K = set(y.astype(int))

        for k in self.K:

            X_k = X[y == k, : ]

            self.likelihoods[k] = {'mean': X_k.mean(axis = 0), "cov": X_k.var(axis = 0) + epsilon}
            self.priors[k] = len(X-k)/len(X)

    def predict(self, X):

        N, D = X.shape

        P_hat = np.zeros((N, len(self.K)))

        for k, l in self.likelihoods.items():
            # Bayes Theorem application
            P_hat[:, k] = mvn.logpdf(X, l['mean'], l['cov']) + np.log(self.priors[k])

        return P_hat.argmax(axis = 1)

# KNN class
class KNNClassifier():

    def fit(self, X, y):
        self.X = X
        self.y = y.astype(int)

    def predict(self, X, k, epsilon = 1e-3):

        N = len(X)
        y_hat = np.zeros(N)
```

```

    for i in range(N):
        dist2 = np.sum((self.X-X[i])**2, axis=1)
        idxt = np.argsort(dist2)[:k]
        gamma_k = 1/(np.sqrt(dist2[idxt] + epsilon))
        y_hat[i] = np.bincount(self.y[idxt], weights = gamma_k).argmax()

    return y_hat

# Gauss-Bayes class
class GaussBayes():

    def fit(self,X,y,epsilon=1e-3):

        self.likelihoods=dict()
        self.priors=dict()
        self.k = set(y.astype(int))

        for k in self.k:
            X_k = X[y==k,:]
            N_k, D = X_k.shape
            mu_k = X_k.mean(axis=0)
            self.likelihoods[k] = {"mean":X.mean(axis=0),"cov":(1/(N_k-1))*np.matmul((X_k-mu_k).T,>
            self.priors[k] = len(X_k) / len(X)

    def predict(self,X):
        N,D = X.shape
        P_hat = np.zeros((N,len(self.k)))
        for k, l in self.likelihoods.items():
            P_hat[:,k] = mvn.logpdf(X,l["mean"],l["cov"])+np.log(self.priors[k])

        return P_hat.argmax(axis=1)

# Accuracy method
def accuracy(y, y_hat):
    return np.mean(y==y_hat)

# min-max normalization method
def min_max_scaling(X):
    new_X =[]
    for i in X:
        new_X.append((i - i.min())/(i.max() - i.min()))
    return np.asarray(new_X)

# Functions for getting and viewing instances
def view_digit(x):
    """ Displays the given instance
    NumpyArray: X -> A row vector with 784 values
    """
    plt.imshow(x.reshape(28,28))

```

```

def get_mean(X, y, k):
    """ returns a row vector with the mean values for the given class
    NumpyArray: X -> An n by m matrix of instances
    NumpyArray: y -> A len(X) row vector of class labels
    Int: K      -> The label in question
    """
    return sum(X[y==k,:]/len(X[y==k,:]))

def view_mean(X, y, k):
    """ Displays the a mean representation of the given class in X
    NumpyArray: X -> An n by m matrix of instances
    NumpyArray: y -> A len(X) row vector of class labels
    Int: K      -> The label in question
    """
    view_digit(get_mean(X, y, k))

```

## ▼ Data Exploration

### Training data

```

# Conversion of the csv files into data frames.
path = '/content/drive/MyDrive/Enhance IT Data Science Course/Week 2/Assignment 2/Data/Mike S
data_train = pd.read_csv(path)
path = '/content/drive/MyDrive/Enhance IT Data Science Course/Week 2/Assignment 2/Data/Mike S
data_test = pd.read_csv(path)

```

```
data_train
```

Unnamed: 0	index	labels	0	1	2	3	4	5	6	...	774	775	776	777	778	7
0	0	0	5	0	0	0	0	0	0	...	0	0	0	0	0	0

▼ Testing data

2	2	2	4	0	0	0	0	0	0	...	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---	---	-----	---	---	---	---	---	---

data\_test

Unnamed: 0	index	labels	0	1	2	3	4	5	6	...	774	775	776	777	778	77
0	0	0	7	0	0	0	0	0	0	...	0	0	0	0	0	0
1	1	1	2	0	0	0	0	0	0	...	0	0	0	0	0	0
2	2	2	1	0	0	0	0	0	0	...	0	0	0	0	0	0
3	3	3	0	0	0	0	0	0	0	...	0	0	0	0	0	0
4	4	4	4	0	0	0	0	0	0	...	0	0	0	0	0	0
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
9995	9995	9995	2	0	0	0	0	0	0	...	0	0	0	0	0	0
9996	9996	9996	3	0	0	0	0	0	0	...	0	0	0	0	0	0
9997	9997	9997	4	0	0	0	0	0	0	...	0	0	0	0	0	0
9998	9998	9998	5	0	0	0	0	0	0	...	0	0	0	0	0	0
9999	9999	9999	6	0	0	0	0	0	0	...	0	0	0	0	0	0



```
X_train = data_train.to_numpy()
```

```
y_train = X_train[:, 2]
```

```
X_train = X_train[:, 3:]
```

```
X_train.shape
```

(60000, 784)

```
y_train.shape
```

(60000,)

```
X_test = data_test.to_numpy()
```

```
y_test = X_test[:, 2]
```

```
X_test = X_test[:, 3:]
```

```
X_test.shape
```

```
(10000, 784)
```

```
y_test.shape
```

```
(10000,)
```

## ▼ Naive Bayes

```
model_nb = GaussNB()
```

```
model_nb.fit(X_train, y_train)
```

```
y_pred_nb_train = model_nb.predict(X_train)
```

```
y_pred_nb_test = model_nb.predict(X_test)
```

```
acc_nb_train = accuracy(y_train, y_pred_nb_train)  
acc_nb_train
```


```
0.5937666666666667
```

```
acc_nb_test = accuracy(y_test, y_pred_nb_test)  
acc_nb_test
```

```
0.5878
```

```
# Confusion matrix  
plt.figure(figsize=(10,7))  
y_actu = pd.Series(y_test, name='Actual')  
y_pred = pd.Series(y_pred_nb_test, name='Predicted')  
cm = pd.crosstab(y_actu, y_pred)  
ax = sns.heatmap(cm, annot=True, fmt="d")  
plt.ylabel('True label')  
plt.xlabel('Predicted label')
```

Text(0.5, 42.0, 'Predicted label')



## ▼ Gauss-Bayes

```

model_gb = GaussBayes()

model_gb.fit(X_train, y_train)

y_pred_gb_train = model_gb.predict(X_train)

y_pred_gb_test = model_gb.predict(X_test)

acc_gb_train = accuracy(y_train, y_pred_gb_train)
acc_gb_train

0.11651666666666667

acc_gb_test = accuracy(y_test, y_pred_gb_test)
acc_gb_test


0.1093

# Confusion matrix

```

```
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_pred_gb_test, name='Predicted')
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

Text(0.5, 42.0, 'Predicted label')



## ▼ KNN

```
model_knn = KNNClassifier()

model_knn.fit(X_train, y_train)

y_pred_knn_train = model_knn.predict(X_train, 12)

y_pred_knn_test = model_knn.predict(X_test, 12)

acc_knn_train = 'Soon!'
acc_knn_test = 'Soon!'
```

```
acc_knn_train = accuracy(y_train, y_pred_knn_train)
acc_knn_train
```

```
1.0
```


```
acc_knn_test = accuracy(y_test, y_pred_knn_test)
acc_knn_test
```

```
0.9678
```

```
# Confusion matrix
```

```
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_pred_knn_test, name='Predicted')
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
Text(0.5, 42.0, 'Predicted label')
```



## Model comparison

```
table = [['Model', 'Train', 'Test'],
         ['Naive Bayes', acc_nb_train, acc_nb_test],
```



```
['Gauss-Bayes', acc_gb_train, acc_gb_test],
['KNN', acc_knn_train, acc_knn_test]]
```

```
print(tabulate(table))
```

```
-----
Model      Train      Test
Naive Bayes 0.5937666666666667 0.5878
Gauss-Bayes 0.11651666666666667 0.1093
KNN         1.0          0.9678
-----
```

## ▼ Data transformation - normalization

```
norm_X_train = min_max_scaling(X_train)
norm_X_test = min_max_scaling(X_test)
```

## ▼ Naive Bayes

```
model_nb = GaussNB()
```

```
model_nb.fit(norm_X_train, y_train)
```

```
y_pred_nb_train = model_nb.predict(norm_X_train)
```

```
y_pred_nb_test = model_nb.predict(norm_X_test)
```

```
acc_nb_train= accuracy(y_train, y_pred_nb_train)
acc_nb_train
```

```
0.7682333333333333
```


```
acc_nb_test= accuracy(y_test, y_pred_nb_test)
acc_nb_test
```

```
0.7746
```

```
# Confusion matrix
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_pred_nb_test, name='Predicted')
```

```
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

Text(0.5, 42.0, 'Predicted label')



## ▼ Gauss-Bayes

```
model_gb = GaussBayes()

model_gb.fit(norm_X_train, y_train)

y_pred_gb_train = model_gb.predict(norm_X_train)

y_pred_gb_test = model_gb.predict(norm_X_test)

acc_gb_train = accuracy(y_train, y_pred_gb_train)
acc_gb_train

0.8657333333333334


acc_gb_test = accuracy(y_test, y_pred_gb_test)
```

```
acc_gb_test
```

```
0.8395
```

```
# Confusion matrix
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_pred_gb_test, name='Predicted')
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
Text(0.5, 42.0, 'Predicted label')
```



## ▼ KNN

```
model_knn = KNNClassifier()
```

```
model_knn.fit(norm_X_train, y_train)
```

```
y_pred_knn_train = model_knn.predict(norm_X_train, 12)
```

```
y_pred_knn_test = model_knn.predict(norm_X_test, 12)
```

```
acc_knn_train = 'Soon!'
acc_knn_test = 'Soon!'
```

```
acc_knn_train = accuracy(y_train, y_pred_knn_train)
acc_knn_train
```

```
1.0
```


```
acc_knn_test = accuracy(y_test, y_pred_knn_test)
acc_knn_test
```

```
0.9678
```

```
# Confusion matrix
```

```
plt.figure(figsize=(10,7))
y_actu = pd.Series(y_test, name='Actual')
y_pred = pd.Series(y_pred_knn_test, name='Predicted')
cm = pd.crosstab(y_actu, y_pred)
ax = sns.heatmap(cm, annot=True, fmt="d")
plt.ylabel('True label')
plt.xlabel('Predicted label')
```

```
Text(0.5, 42.0, 'Predicted label')
```



## ▼ Model comparison

```
table = [['Model', 'Train', 'Test'],  
         ['Naive Bayes', acc_nb_train, acc_nb_test],  
         ['Gauss-Bayes', acc_gb_train, acc_gb_test],  
         ['KNN', acc_knn_train, acc_knn_test]]
```

```
print(tabulate(table))
```

Model	Train	Test
Naive Bayes	0.7682333333333333	0.7746
Gauss-Bayes	0.8657333333333334	0.8395
KNN	1.0	0.9678

✓ 0s completed at 7:03 PM

